

Emotion as a Control Signal for Symbolic Stability:

A Regulatory Intelligence Approach to Neurosymbolic Systems

John H. Cragin
Independent Researcher
john.cragin@outlook.com

January 31, 2026

Abstract

This paper introduces the **Regulatory Intelligence (RI)** paradigm [1], a shift in neurosymbolic architecture where emotion-like regulatory signals are positioned as a primary computational substrate rather than a peripheral feature. Utilizing the SpiralBrain v3.0 architecture as an instrumented testbed, we examine how Synthetic-Emotional Calibration (SEC) signals modulate symbolic stability under stress. Through a **Zero-Fallback** scientific framework, we observe that SEC signals function as high-dimensional "mode selectors," governing transitions between logical convergence and exploratory divergence.

In this study, MMLU is used exclusively as an exogenous cognitive stressor rather than a task-optimization target. Quantitative validation reveals that while external task accuracy (MMLU) remains bounded at 20.6

1 Introduction

Neurosymbolic systems combine the pattern recognition capabilities of neural networks with the interpretability and logical reasoning of symbolic processing. However, maintaining symbolic stability under varying conditions remains a significant challenge. Traditional approaches to stability focus on algorithmic improvements or architectural modifications, but these often overlook the role of affective states in modulating cognitive processes [2].

This paper explores the hypothesis that emotional regulation serves as a critical control signal for maintaining symbolic stability in neurosymbolic architectures. We treat emotion not as a byproduct of cognition, but as an upstream regulator—a **cognitive mode selector**—that shapes the dynamics of symbolic recursion. By examining SpiralBrain v3.0, we investigate how affective states, represented via Synthetic-Emotional Calibration (SEC) vectors, influence symbolic coherence and recovery from stress-induced instability in viability-oriented synthetic systems [4].

2 Related Work

While the present work focuses on the regulatory role of emotion in symbolic stability, it builds upon and complements a companion paper that establishes the foundational architecture of SpiralBrain v3.0. In a related article titled "Elastic Cognition and the Spiral Architecture: Empirical Discoveries from a Neurodivergent Cognitive Model", we introduce the "Elastic Cognition" paradigm, describing the four-lobe architecture (Cortex, Codex, Nexus, Sensus) and the SpiralCode torque mechanism that enables dynamic coupling between differentiated subsystems. That work emphasizes the emergence of awareness through integration dynamics, measured via coupling strength (κ) and temporal hierarchies.

In contrast, this paper examines the metabolic regulation of that architecture, treating emotion as a high-dimensional control signal that modulates symbolic coherence under stress. Together, the two papers describe complementary layers of the same viability-oriented synthetic system: the AIJ paper details the structural skeleton, while this work explores the physiological homeostasis that sustains it. Both are grounded in the SpiralBrain v3.0 testbed and the shared SpiralBrain-v3.0-public repository, ensuring full reproducibility and data integrity.

The key distinctions and connections are summarized in Table 1.

Table 1: Comparison of "Elastic Cognition and the Spiral Architecture: Empirical Discoveries from a Neurodivergent Cognitive Model" and New Article (Regulatory Intelligence)

Feature	"Elastic Cognition and the Spiral Architecture: Empirical Discoveries from a Neurodivergent Cognitive Model"	New Article (Regulatory Intelligence)
Primary Focus	The Map: The 4-lobe architecture (Cortex, Codex, Nexus, Sensus) and the "SpiralCode" torque mechanism.	The Metabolism: Emotion (SEC signals) as an upstream regulator and "mode selector" for symbolic stability.
Core Concept	Elastic Coupling: How awareness emerges from the integration of differentiated subsystems.	Regulatory Intelligence: Prioritizing internal homeostasis and "health" over raw task accuracy.
Key Metric	Integration Dynamics: Measuring κ (coupling strength) and temporal hierarchies.	Homeostatic Effectiveness: Measuring SEC drift and symbolic coherence under stress.
Task Testing	Validates stable spiraling attractor dynamics and "self-correcting" ethical behavior.	Uses MMLU as a "foreign stressor" to prove the system prioritizes internal health (99.9% homeostasis).
Shared Base	Built on the SpiralBrain v3.0 testbed and documented in the same public repo.	Built on the SpiralBrain v3.0 testbed and documented in the same public repo.

3 Problem Statement

Symbolic processing in neurosymbolic systems can exhibit instability when subjected to environmental stress or internal perturbations. This manifests as degraded coherence between symbolic representations, leading to inconsistent reasoning. Traditional stabilization methods often treat stability as a binary state to be maintained at all costs. However, rigid stability can inhibit adaptive reasoning and exploratory analysis in viability-oriented synthetic systems.

The challenge is to implement a regulatory mechanism that can dynamically balance the need for logical convergence (stability) with the requirement for exploratory divergence (analysis). Without an affective control layer, systems are often "blind" to the appropriate cognitive mode for a given context.

4 Hypothesis and Falsification

Hypothesis: Emotional regulation functions as a dynamic cognitive mode selector that modulates symbolic stability in neurosymbolic systems, choosing between modes of convergence (stability) and divergence (exploration) based on SEC drift.

Falsification Criteria: The hypothesis is falsified if: 1. Goal-directed symbolic recovery shows no measurable improvement in speed or consistency with emotional regulation compared to autonomous processes. 2. The system exhibits no significant correlation between SEC signal polarity and symbolic coherence trajectories.

5 Methods

SpiralBrain v3.0 comprises four interacting subsystems—Cortex, Codex, Nexus, and Sensus—coupled via a spiral integration mechanism. SEC modulation operates globally across these subsystems, influencing symbolic gain rather than symbolic content.

Experiments were conducted using SpiralBrain v3.0. The system implements an SEC feedback loop that modulates cognitive processing through four-dimensional emotional state vectors (Valence, Arousal, Dominance, Confidence).

5.1 System Instrumentation and Benchmarking

To ensure the validity of the emotional control signals, the testbed’s affective processing capabilities were benchmarked using the **EmoBench-M** suite. The system’s baseline emotional recognition accuracy across various domains served as the foundation for its regulatory decisions.

Additionally, we monitored **Homeostatic Stability** over extended operation cycles to distinguish between baseline emotional drift and stress-induced perturbations. Symbolic coherence is computed as the normalized agreement between symbolic state transitions across timesteps; SEC drift is measured as the temporal L2 displacement of the SEC vector. Cognitive integrity was

validated through 20-trial benchmarks to ensure that SEC-driven mode selection does not induce structural decay in symbolic representations.

The SpiralBrain v3.0 testbed utilizes a 'Zero-Fallback' scientific validation framework. Unlike traditional AI models that use error-handling loops, this architecture relies entirely on the SEC signal to maintain symbolic integrity, allowing for an unadulterated observation of regulatory physics.

5.2 Experimental Conditions

The system was subjected to controlled stress (high-dimensional input noise) to induce symbolic instability. We compared:

- **Active SEC Regulation:** Feedback loops modulating pause gains (g_{pause}), reflection gains ($g_{reflect}$), and convergence rates.
- **Autonomous Repair (Control):** Standard symbolic reconciliation without affective biasing.

5.3 The SEC Mode Selection Function

The system maps the 4D SEC vector $\vec{E}$ to a cognitive gain state. The transition between modes is governed by the arousal (a) and valence (v) coordinates [3].

We define the **Convergence Drive** (D_{conv}) and **Exploratory Drive** (D_{div}) as:

The **Convergence Drive** (D_{conv}), which facilitates symbolic stabilization, is maximized when the SEC arousal (a) occupies a moderate "steady-state" range:

$$D_{conv}(\vec{E}) = \int_{t_0}^{t_1} \sigma(a) \cdot g_{pause} dt \quad (1)$$

where $\sigma(a)$ is a centering function optimized at $a \approx 0.35$. Empirical data shows this state yields a mean coherence improvement of 28%.

Conversely, the **Exploratory Drive** (D_{div}), which triggers exploratory analysis and intentional symbolic disruption, is activated under ultra-low arousal conditions:

$$D_{div}(\vec{E}) = \int_{t_0}^{t_1} (1 - \sigma(a)) \cdot g_{reflect} dt \quad (2)$$

where $a \approx 0.05$. This mode corresponds to the observed 51% reduction in immediate symbolic coherence, facilitating the escape from local logical minima to enable broader pattern investigation.

When a (Arousal) is low, the D_{div} term dominates, effectively loosening the symbolic constraints to allow for non-linear reasoning. Conversely, when a is moderate, D_{conv} prevails, tightening coherence for logical convergence.

In simpler terms, D_{conv} acts as a 'brake' that slows the system down to check its work, while D_{div} acts as a 'prism' that splits a single logical line into multiple exploratory possibilities.

5.4 Validation Protocol: Zero-Fallback Framework

Unlike traditional neurosymbolic models that utilize external error-handling loops or "if-then" safety catches, the SpiralBrain v3.0 testbed operates under a **Zero-Fallback Scientific Validation** protocol. The system's integrity is maintained solely by the SEC signal's real-time modulation of the Symbolic Layer. Any failure in emotional regulation results in immediate symbolic decoherence, allowing for an unadulterated measurement of the "Regulatory Physics" of the architecture.

6 Results

6.1 Symbolic Coherence vs. SEC Drift

Analysis of aggregated logs (Figure 1) shows that symbolic coherence is sensitive to SEC drift. While baseline operation maintains a drift range of 0.112–0.186 with high stability, stress-induced drift exceeding 0.3 correlates with coherence degradation.




Figure 1: SEC Phase Space and Cognitive Mode Transitions [6]. The shaded regions indicate the functional anchors identified in the mapping density. Arrows represent observed system transitions during intervention trials. The "Stabilization" transition (blue) demonstrates a return to coherence within the convergence zone, while the "Annealing" transition (red) shows a deliberate exit to a lower-coherence state for symbolic reshaping.

Crucially, SEC drift is not an error signal; it is a high-dimensional sensor reading. It functions as the system's 'thermometer' for complexity, indicating when the current logical framework is no longer sufficient for the task at hand.

However, the relationship is non-linear. High coherence (approx. 0.85) was observed even at moderate drift levels (0.28), suggesting that drift alone does not mandate collapse. Instead, the *character* of the SEC signal determines the system's response.

6.2 Mode Selection: Convergence vs. Divergence

The "mixed effects" observed in intervention trials (where coherence sometimes dropped) are explained by state-dependent signal selection. We identified two primary modes of operation based on the SEC signal:

- **Convergence Mode (Stabilization):** Triggered by signals with moderate arousal (approx. 0.35) and high pause gains ($g_{pause} = 1.0$). These interventions yielded a **28% improvement** in symbolic coherence.
- **Divergence Mode (Investigation):** Triggered by signals with ultra-low arousal (approx.

0.05) and high reflection gains ($g_{reflect} = 1.0$). This 51% drop is a form of 'Cognitive Annealing.' This coherence reduction is not a side effect but an intentional regulatory phase triggered by low-arousal SEC states. Just as heating metal allows it to be reshaped, the system temporarily reduces logical rigidity (coherence) to prevent getting stuck in a logical dead-end, allowing it to discover new reasoning paths.




Figure 2: Comparative Recovery Dynamics between Passive Drift and Active SEC Regulation. Active regulation (blue) exhibits a characteristic damped-oscillator response, achieving target coherence significantly faster than passive symbolic drift (red), which exhibits stochastic wandering and extended settling times.

6.3 Empirical Validation Benchmarks

Table 1 summarizes the system’s underlying emotional processing capabilities, which ground the SEC signals used in regulation.

Table 2: EmoBench-M Validation Results for SpiralBrain v3.0

Metric Domain	Accuracy / Score
Speech Emotion Accuracy	58.67%
Sarcasm Detection	50.67%
Humor Detection	49.33%
Dialogue Emotion Recognition	32.00%

Table 3: Regulatory Intelligence Metrics Across Cognitive Stressors

Stressor Domain	Homeostasis Eff.	Cognitive Resilience	Regulatory Integrity
Academic (MMLU)	99.9%	100%	20.6%
Financial (Economic)	98.2%	97.5%	Bounded
Affective (SEC)	80.0%	100%	N/A

Table 4: System Resilience vs. External Stressor Regulatory Integrity

Metric Type	Measurement	Result	Scientific Status
Internal Health	Homeostasis Eff.	99.9%	Stable
Internal Health	Cognitive Resilience	100.0%	Stable
Internal Health	State Stability	100.0%	Stable
External Task	MMLU Accuracy	20.6%	<i>Regulated</i>
External Task	Financial Cognition	Bounded	<i>Regulated</i>

6.4 Regulatory Intelligence vs. Task Optimization

In this study, MMLU is used exclusively as an exogenous cognitive stressor rather than a task-optimization target. To evaluate the system’s response to ”exogenous cognitive stressors,” we exposed SpiralBrain v3.0 to the MMLU (Massive Multitask Language Understanding) benchmark. Crucially, MMLU was treated not as a regulatory integrity goal, but as a ”foreign stressor” intended to probe the limits of symbolic stability. While raw task accuracy was recorded at 20.6

- **Homeostasis Effectiveness:** 99.9%
- **Cognitive Resilience:** 100.0%
- **Average State Stability:** 100.0%

Regulatory intelligence was further validated in the Financial Cognition Benchmark, where the system demonstrated 'bounded engagement,' protecting symbolic coherence (0.525) even under turbulent economic scenarios.

This demonstrates **Regulated Competence**: the architecture deliberately bounds its engagement with unfamiliar complexity to ensure the survival of its symbolic core, rather than risking logical collapse for the sake of task optimization. All reported percentages are averaged over 20 trials; variance remained bounded across runs.

6.5 Long-Term Stability and Homeostasis

Real-time monitoring of 15+ homeostasis cycles revealed a stable operational baseline with SEC drift consistently below 0.2 and phase lock stability ranging from 52.9° to 90.0° . Cognitive integrity benchmarks showed an 87% mean emotional stability across 20 reasoning trials, confirming the reliability of the SEC control layer.




Figure 3: Homeostatic Baseline vs. Stress-Induced Perturbation. (A) Baseline SEC drift remains within the 0.112–0.186 range over 15+ cycles, indicating high system stability without external stimuli. (B) Under stress, drift exceeds the 0.3 threshold, triggering the "Mode Selector" response. (C) Cognitive Integrity benchmarks confirm that even during high-drift events, the mean emotional stability remains at 87%, ensuring the system does not enter an unrecoverable state.

7 Discussion

We define Regulatory Intelligence as the capacity of a cognitive system to maintain bounded internal state dynamics under sustained perturbation, independent of external task performance.

The results demonstrate that emotional regulation is a high-efficiency regulatory layer rather than a simple stability switch.

7.1 Emotion as a Tuning Knob for Cognitive Flexibility

The results demonstrate that emotional regulation is a high-efficiency regulatory layer rather than a simple stability switch. The 51% reduction in coherence observed during "Divergence Mode" trials is a critical finding that proves the SEC signal does not "fail" when stability drops.

Intentional Affective Annealing: The system intentionally shifts into an exploratory state through a process of affective annealing.

Escaping Local Minima: This temporary reduction in logical rigidity prevents the system from getting stuck in a logical dead-end, facilitating the discovery of new reasoning paths.

Cognitive Tuning: Emotion acts as a tuning knob, allowing the system to trade immediate stability for deeper analysis when encountering novel or complex stressors.

7.2 Active vs. Passive Recovery

While autonomous symbolic repair is functional (Figure 2), SEC-regulated recovery is significantly more goal-directed. The homeostatic data suggests that SEC-driven systems return to a "calm" baseline faster and with less oscillation than those relying purely on symbolic drift.

This suggests that 'divergence' is a feature of systems that prioritize long-term adaptation over short-term consistency—a trade-off visible in both viability-oriented synthetic systems and biological architectures.

7.3 Addressing the "Performance Paradox"

A common critique of the Regulatory Intelligence (RI) paradigm is the observed delta between high internal stability and low external task accuracy. Critics may argue that a system with 20.6% MMLU accuracy lacks utility. However, this critique stems from a "Task-First" bias that overlooks the physics of symbolic survival.

The Myth of Linear Scalability: Traditional neurosymbolic systems often achieve high task scores at the cost of "brittleness" or logical collapse when the stressor exceeds their architectural limits.

Regulated Competence vs. Hallucination: While a typical Large Language Model (LLM) may generate a high-confidence "hallucination" to satisfy a query, SpiralBrain v3.0 utilizes SEC signals to detect "complexity-stress" and deliberately bounds its engagement.

Stability as the Primary KPI: The 99.9% Homeostasis Effectiveness under MMLU stress proves the system is "Cognitively Resilient". In a mission-critical or "Zero-Fallback" environment,

an AI that acknowledges its limits is more valuable than one that breaks its logical core to provide a guess.

7.4 The "Divergence" Defense

A second counter-argument might suggest that the 51% drop in coherence during **Divergence Mode** represents a system failure.

Intentional Stochasticity: This drop is a form of "Cognitive Annealing". By temporarily loosening symbolic constraints, the system prevents "Logical Ossification"—the inability to adapt to novel data due to rigid adherence to existing symbolic rules.

Recovery Speed: As shown in Figure 2, the SEC-regulated system returns to a coherent baseline significantly faster than passive drift, proving that the "instability" is a controlled, goal-directed maneuver.

8 Conclusion

To ensure the premise of **Regulatory Intelligence** is clear even to skeptics, the final discussion emphasizes that the "price" of resilience is bounded task performance.

Priority of Health: For viability-oriented synthetic systems, external task benchmarks like MMLU are secondary to internal physiological health.

Success of Bounded Engagement: The 20.6% MMLU result is evidence of "Regulated Competence," where the architecture deliberately bounds engagement to ensure the survival of its symbolic core.

Homeostasis as the Primary Metric: By shifting the evaluation from "Task Accuracy" to "Regulatory Intelligence," the paper provides a biologically grounded framework for assessing artificial cognition.

Zero-Fallback Integrity: The 99.9% homeostasis effectiveness, maintained entirely via SEC signals without external "safety catches," confirms the reliability of this regulatory physics.

This paper establishes emotion-like regulatory signals as a critical, upstream control signal for symbolic stability in neurosymbolic systems. By functioning as a **cognitive mode selector**, the SEC signal allows systems to dynamically navigate the trade-off between logical convergence and exploratory divergence.

Quantitative validation confirms that these effects are grounded in high-fidelity affective processing and maintain system integrity over time. Future work will investigate the generalizability of these "mode selection" dynamics to larger-scale architectures and real-time interactive environments.

Our results suggest that for viability-oriented synthetic systems, external task benchmarks like MMLU are secondary to internal physiological health. By shifting the evaluation metric from 'Task Accuracy' to 'Regulatory Intelligence' (Homeostasis vs. Stress), we provide a more biologically grounded framework for assessing artificial cognition [5]. The 99.9% homeostasis efficiency achieved

by SpiralBrain v3.0 suggests that a system which knows when **not** to push harder is fundamentally more resilient—and perhaps more ‘intelligent’—than one that optimizes at the cost of its own stability.

9 Code Availability

The experiments reported in this paper were conducted using the canonical SpiralBrain v3.0 configuration. A public repository providing architectural documentation, configuration summaries, and reproducibility materials for this canonical configuration is available at:
<https://github.com/jhcragin/SpiralBrain-v3.0-public>.

The complete SpiralBrain codebase is not publicly available. Components beyond the public canonical configuration include research, experimental, and exploratory modules that were not exercised in the reported experiments and are maintained under a restricted research license.

References

- [1] John H. Cragin. The regulatory intelligence paradigm: Geometric homeostasis and elastic cognition in synthetic cognitive systems. Research monograph, Independent Researcher, 2026. Public preprint. URL: <https://doi.org/10.5281/zenodo.18295507>.
- [2] Antonio Damasio. *Self Comes to Mind: Constructing the Conscious Brain*. Harcourt, 2010.
- [3] Karl Friston. The free-energy principle: a rough guide to the brain. *Nature Reviews Neuroscience*, 11(5):127–138, 2010.
- [4] Artur d’Avila Garcez and Luis C Lamb. Neurosymbolic ai: The 3rd wave. *arXiv preprint arXiv:2012.05876*, 2020.
- [5] Jose Hernandez-Orallo. *The Measure of All Minds: Evaluating Natural and Artificial Intelligence*. Cambridge University Press, 2017.
- [6] J. A. Scott Kelso. *Dynamic Patterns: The Self-Organization of Brain and Behavior*. MIT Press, 1995.