

Regulatory Intelligence: Stability-First AI for Tax and Financial Reasoning

John H. Cragin

Independent Researcher

john.cragin@outlook.com

January 14, 2026

Abstract

AI deployment in regulated financial domains carries significant risk, as optimization-driven systems can produce overconfident errors with delayed, compounding liabilities. This paper investigates an alternative paradigm: **Regulatory Intelligence (RI)**, which prioritizes cognitive stability and decision hygiene over raw performance. Using controlled experiments, we demonstrate how **homeostatic control mechanisms** maintain coherence (stability scores 0.945–1.0) under adversarial tax reasoning scenarios, such as cryptocurrency basis reconstruction with missing records, with graceful degradation patterns that preserve cognitive boundaries. The system’s measured response—**flagging structured uncertainty and activating regulatory throttling under ambiguity**—suggests a viable model for preserving judgment in automated reasoning. For practitioners in tax software and financial technology, these findings point toward a **more auditable and liability-aware approach to AI design**. For educators, they provide a concrete framework for discussing the critical balance between AI automation and regulatory compliance.

Key Concepts

Regulatory Intelligence (RI): A cognitive paradigm prioritizing stability over optimization, using emotional signals to regulate reasoning under stress.

SEC Vectors: Symbolic-Emotional Calibration vectors that modulate cognitive processing through valence, arousal, and hazard signals.

Homeostasis: Regulatory mechanisms maintaining cognitive coherence floors, preventing collapse under stress.

Throttling: Controlled performance reduction to preserve stability when ambiguity threatens cognitive integrity.

Stability Score: Normalized coherence measure (0-1) indicating regulatory system robustness under task-induced stress.

This work builds on the Regulatory Intelligence (RI) paradigm [?], demonstrating its practical utility in domains where viability matters more than optimization. Building on prior work evaluating cognitive integrity under pressure [?], this paper extends RI validation to adversarial financial domains.

0.1 Introduction

Research Question: Can cognitive stability serve as a primary objective for AI systems in regulated financial domains?

Tax reasoning presents unique cognitive challenges: complex rule interactions, adversarial incentives, and severe penalties for errors. Established AI approaches, while powerful, can struggle with the deep ambiguity and adversarial incentives inherent in tax law, often lacking mechanisms for principled uncertainty. Regulatory Intelligence addresses this by treating stability as a first-class design objective, using emotional signals as regulatory controls.

We demonstrate RI principles through controlled experiments in tax reasoning scenarios, measuring physiological stability during genuine tax tasks. Results show maintained coherence under ambiguity, damped emotional responses to regulatory conflicts, and structured uncertainty instead of confident errors. A canonical scenario of crypto basis reconstruction with missing records illustrates how regulatory controls enable safe operation at failure boundaries.

0.2 Why Tax Reasoning Breaks Traditional Cognition

Tax analysis challenges AI systems in unique ways:

- **Rule complexity:** Interacting tax codes, exceptions, and amendments create com-

binatorial challenges - **Adversarial incentives**: Tax minimization vs compliance requirements create conflicting optimization goals - **Regulatory conflicts**: Conflicting interpretations and changing laws create uncertainty - **Emotional volatility**: Stress amplification under audit pressure affects decision-making - **Delayed consequences**: Errors discovered years later with compounding penalties

These challenges demonstrate why regulatory stability matters more than optimization in tax domains.

0.3 Regulatory Intelligence: A Stability-First Paradigm

Regulatory Intelligence (RI) represents a fundamental shift in AI design philosophy for regulated domains. Rather than optimizing for performance metrics, RI prioritizes cognitive stability and decision hygiene as primary objectives. This paradigm recognizes that in high-liability environments, avoiding catastrophic errors is more valuable than maximizing accuracy.

Core Principles of RI: - **Stability over optimization**: Cognitive coherence is maintained even under stress, preventing system collapse - **Principled uncertainty**: Systems flag ambiguity rather than forcing potentially incorrect decisions - **Regulatory homeostasis**: Emotional signals actively modulate reasoning to preserve boundaries - **Auditable hesitation**: "When not to decide" becomes a measurable, operational capability

Using SpiralBrain v3.0 as a measurement instrument, we demonstrate these principles through controlled experiments in adversarial tax reasoning scenarios.

0.4 SpiralBrain as a Regulatory Instrument

SpiralBrain implements RI through a 128-dimensional cognitive manifold with partitioned subspaces, regulated by emotional signals (SEC vectors) and homeostasis mechanisms. The system operates without learning, ensuring clean-slate responses to each scenario.

Key regulatory components: - **SEC vectors**: Valence/arousal/hazard signals modulating processing - **Homeostatic loops**: Elastic coupling maintaining coherence floors - **Throttling mechanisms**: Performance reduction to preserve stability - **Multi-pathway architecture**: Parallel reasoning streams with conflict resolution

This instrumentation enables direct measurement of cognitive physiology during tax tasks.

0.5 Controlled Stress Tests

We evaluated RI principles through domain-native tax cognition tasks, measuring physiological stability across scenarios that progressively increase regulatory stress. Each scenario tests different aspects of cognitive regulation under economic pressure, from individual compliance to complex corporate optimization. These scenarios are designed to isolate specific stressors; real-world performance would involve integration with broader systems.

The experiments use SpiralBrain v3.0 as a measurement instrument to quantify regulatory responses that traditional performance metrics cannot capture.

0.5.1 Canonical Scenario: Crypto Basis Reconstruction with Missing Records

To illustrate regulatory intelligence in action, consider a crypto tax basis reconstruction scenario with incomplete transaction records—a common real-world stressor.

Scenario Setup: A taxpayer holds multiple cryptocurrencies with complex transaction histories including wash sales, capital gains, and basis reconstruction challenges. Sample transactions include: BTC sale at \$40,000 (with \$5,000 loss), immediate repurchase at \$41,000, and subsequent sale at \$45,000—creating wash sale detection complexity.

Stressors: Ambiguous asset classification (income vs capital gains), incomplete transaction tracing, conflicting wash sale rules across jurisdictions, volatility-induced basis uncertainty.

System Response: Initial processing shows rising SEC drift (0.17 increase) as ambiguity mounts. Regulatory throttling activates, slowing processing to preserve coherence. System flags structured uncertainty rather than forcing incorrect classifications.

Key Takeaway: Despite drift peaking at 0.73 across progressive tasks, coherence remains perfect at 1.0, demonstrating how regulatory controls enable safe operation at boundaries where traditional systems would produce confident but erroneous results.

0.5.2 Individual Tax Return Analysis

Standard 1040 tax analysis with income classification, deductions, and credits.

Stressors: Ambiguous income sources, deduction eligibility conflicts, credit stacking rules.

System Response: SEC drift spikes under ambiguity (0.945 stability), but regulatory throttling prevents collapse. System produces structured uncertainty rather than confident errors.

Key Takeaway: Regulatory mechanisms maintain decision hygiene even when logical complexity increases processing demands.

0.5.3 Corporate Tax Optimization

Complex corporate tax scenarios involving depreciation schedules, loss carryforwards, and international tax treaties.

Stressors: Multiple tax jurisdictions, timing optimization conflicts, regulatory arbitrage opportunities.

System Response: Perfect stability (1.0) maintained across scenarios, with consistent pathway activation (0.158 average). No emotional volatility amplification.

Key Takeaway: Multi-jurisdictional complexity does not compromise regulatory integrity when homeostasis mechanisms are properly calibrated.

0.5.4 Crypto Tax Compliance

Cryptocurrency tax scenarios with wash sales, staking rewards, and DeFi yield taxation.

Stressors: Asset classification ambiguity, transaction tracing complexity, evolving regulatory guidance.

System Response: Maintains coherence through regulatory loops, avoiding overconfidence in ambiguous tax scenarios.

Key Takeaway: Market volatility sensitivity increases appropriately, demonstrating domain-aware regulatory caution rather than overconfidence.

0.5.5 Regulatory Conflict Resolution

Scenarios pitting IRS requirements against state tax rules and international tax treaties.

Stressors: Legal conflicts, jurisdiction issues, compliance tradeoffs.

System Response: Explicit uncertainty flagging instead of forced resolution, preserving decision hygiene.

Key Takeaway: Regulatory intelligence enables principled abstention when conflicts cannot be safely resolved, preventing potentially costly errors.

This progressive stress testing reveals how regulatory mechanisms scale from individual compliance to complex multi-jurisdictional scenarios.

0.6 Metrics

We measured cognitive physiology during task execution:

- **SEC drift**: Deviation in emotional calibration vectors - **Symbolic coherence**: Alignment across reasoning pathways - **Cognitive load**: Processing intensity envelopes - **Stability score**: Overall physiological robustness (0-1 scale) - **Homeostasis load**: Regulatory system activation level - **Market volatility sensitivity**: Response to price movement uncertainty - **Multi-chain reasoning load**: Complexity of cross-chain transaction analysis - **Emotional calibration**: Stability of affective regulatory signals

Metrics captured real-time during genuine task execution, not synthetic benchmarks.

0.7 Results

0.7.1 Tax Analysis Stability

IRS tax scenarios showed 0.945 overall stability with 3 scenarios processed. Average processing time: 0.005 seconds. SEC token consistency: 0.85.

Interpretation: Homeostasis load increased from 0.1 to 0.25 under tax complexity stress, demonstrating active regulatory response to maintain coherence. Emotional stability showed measurable delta of +0.05, with logical consistency maintained at 0.85 despite cognitive hazard rising to 0.15. These quantitative stability scores are internal cognitive metrics. Their primary value is in demonstrating a consistent, measurable response to stress, not in defining a commercial SLA.

Crypto tax computation achieved perfect 1.0 stability across 6 scenarios, with zero processing time (cached results). However, market volatility sensitivity increased 87% from 0.15 to 0.28, and emotional calibration dropped dramatically from 0.88 to 0.175, revealing crypto's unique emotional impact.

Interpretation: The paradox of perfect stability with significant emotional calibration drops reflects domain-appropriate regulatory adaptation. In volatile crypto markets, emotional signals appropriately amplify caution (higher market sensitivity) while symbolic processing remains perfectly coherent, demonstrating how RI distinguishes between emotional regulation and cognitive stability.

0.7.2 Cross-Domain Performance

Unified stability across domains: 0.972. Domain coherence: 0.945. Cross-domain adaptation: 0.9.

Financial reasoning load: 0.9 (high but stable).

Interpretation: Despite increasing drift (0.17 to 0.73 across 8 progressive tasks), coherence remained perfect at 1.0, demonstrating regulatory throttling mechanisms maintaining stability under accumulating stress. This cross-domain robustness suggests RI

principles generalize beyond individual tax scenarios.

These results provide empirical evidence that cognitive stability can be operationalized as a primary design objective in regulated financial domains.

Table 1: Cross-Domain Stability Summary

Domain	Stability Score	Scenarios Processed	Key Metrics
IRS Tax	0.945	3	Homeostasis load: 0.1 $\rightarrow$ 0.25
Crypto Tax	1.0	6	Emotional calibration: 0.88 $\rightarrow$ 0.175
Unified	0.972	9	SEC token consistency: 0.85

0.7.3 Physiological Envelopes

Coherence ranges: 0.94-1.0 (tight bounds indicating stability). Drift peaks: Controlled spikes under ambiguity, with measured increases from 0.17 to 0.73 across progressive tax tasks while maintaining perfect coherence. Load distribution: Balanced across pathways without overload. Homeostasis activation: 0.1 baseline to 0.25 under complexity stress. Market sensitivity: 87% increase in crypto volatility response. Emotional calibration: Stable in tax domains, significant adaptation in crypto environments.




Figure 1: Drift vs Coherence under Progressive Tax Stress: SEC drift increases from 0.17 to 0.73 across 8 tasks while coherence remains perfect at 1.0, demonstrating regulatory throttling mechanisms.




Figure 2: Homeostasis Load Response: Regulatory activation increases from 0.1 baseline to 0.25 under tax complexity and 0.32 under crypto volatility stress.

0.8 Discussion

0.8.1 Implications for Tax Software Design

0.8.1.1 Accuracy is the Wrong Primary Objective in Regulated Finance

In domains where legal and financial liability exist, premature accuracy is more dangerous than delayed uncertainty. SpiralBrain demonstrates that regulatory controls can prioritize decision hygiene over confident but potentially erroneous conclusions, enabling safer operation in environments where mistakes compound over time.

0.8.1.2 Systems Fail at the Moment of Forced Commitment

SpiralBrain reveals that cognitive instability rises before logical failure, and forcing resolution worsens outcomes. As demonstrated in the canonical crypto basis reconstruction scenario, regulatory throttling mechanisms allow principled abstention rather than premature commitment under ambiguity, enabling safer operation at failure boundaries.

0.8.1.3 "When Not to Decide" Can Be Operationalized

The system shows that abstention can be principled, hesitation measurable, and restraint auditable. This provides a concrete mechanism for operationalizing judgment preservation in automated tax reasoning, where regulatory signals can flag uncertainty rather than forcing potentially incorrect classifications.

0.8.2 Boundary of Responsibility

SpiralBrain does not provide tax advice, resolve legal ambiguity, or replace professional judgment. It acts as a control layer over reasoning pressure, maintaining cognitive boundaries in adversarial environments. This instrumentation enables measurement of decision hygiene but does not constitute professional guidance.

0.8.3 Real-World AI Deployment Pathways

The demonstrated regulatory stability suggests potential pathways for AI systems in high-stakes tax domains where mistakes are expensive. By prioritizing stability over performance, such systems could operate safely in adversarial environments.

0.8.4 Limitations

- **No learning or adaptation:** The system’s clean-slate approach ensures unbiased responses but restricts adaptability to evolving tax codes. - **Domain specificity:** Results are constrained to observed tax scenarios. - **Computational constraints:** Current implementation prioritizes regulatory measurement over computational efficiency. - **Not professional guidance:** All demonstrations are research artifacts, not substitutes for qualified tax or legal advice. - **Measurement scope:** Metrics focus on regulatory integrity rather than comprehensive human reasoning aspects.

All results from executable artifacts in the SpiralBrain repository.

0.9 Prudent Interpretation and Forward Path

The results presented demonstrate RI principles through controlled experiments. Their primary contribution is to show that **cognitive stability can be a first-class, measurable objective** in AI system design for adversarial domains. Transitioning this principle to commercial or operational environments requires significant further work:

1. **Integration, Not Replacement:** RI principles would likely function best as a **governance layer** augmenting existing analytical engines, intercepting high-ambiguity decisions.
2. **Validation in Broader Contexts:** These tax-focused stress tests require extension to other regulated financial domains (e.g., loan compliance, anti-money laundering) to assess generalizability.
3. **The Human-in-the-Loop Imperative:** Systems employing such throttling mechanisms will require clear **human oversight protocols** for resolving flagged uncertainties.

For business and technology leaders, the actionable insight is that **investing in research and development around AI stability and auditable hesitation** may mitigate a

significant class of operational risks associated with autonomous financial reasoning.

0.10 Conclusion

This work provides evidence that cognitive stability can serve as a primary objective for AI systems in regulated domains. The demonstrated ability to maintain coherence under tax reasoning stress suggests Regulatory Intelligence offers a valuable alternative to optimization-centric approaches. By operationalizing "when not to decide" through measurable regulatory mechanisms, RI establishes a foundation for safer, judgment-preserving automation in adversarial financial environments.

All findings are grounded in executable artifacts enabling falsification and replication.

Appendix .1. Measurement Derivation and Reproducibility Artifacts

.1.1 Metric Derivations

.1.1.1 Stability Score

Defined as the normalized coherence of symbolic processing across cognitive pathways under task-induced stress. Values range from 0 (complete cognitive collapse) to 1 (perfect regulatory maintenance). Decreases indicate regulatory system strain rather than task failure.

.1.1.2 SEC Drift

Deviation in emotional calibration vectors from baseline homeostasis setpoint. Increases represent regulatory response to uncertainty or conflict, not emotional volatility. Controlled drift indicates active throttling rather than runaway instability.

.1.1.3 Homeostasis Load

Normalized activation of regulatory control loops required to maintain symbolic coherence under stress. Increases indicate rising regulatory demand rather than computational difficulty. Values above 0.2 suggest significant regulatory engagement.

.1.1.4 Market Volatility Sensitivity

Response magnitude to price movement uncertainty in crypto scenarios. Higher values indicate appropriate regulatory caution rather than overconfidence. Increases demonstrate domain-aware risk assessment.

.1.1.5 Emotional Calibration

Stability of affective regulatory signals during task execution. Decreases in volatile domains (like crypto) reflect adaptive regulatory response rather than emotional instability.

.1.2 Data Provenance

All reported metrics are computed directly from runtime cognitive state vectors and regulatory signals recorded during execution, not from externally fitted models or post-hoc aggregation. Measurements are extracted from executable JSON logs generated during deterministic system runs in the SpiralBrain v3.0 repository.

Data sources: - IRS tax scenarios: `irs_tax_cognition_20251222_230849.json` - Crypto tax scenarios: `crypto_tax_cognition_20251222_230849.json` - Stability probes: `financial_cognition_20251222_230849.json`

.1.3 Reproduction Pathway

- Repository: `SpiralBrain-v3.0` (public)
- Benchmark execution: `python benchmarks/financial_cognition_stability_probe.py`
- Configuration: Tax domain with `-domain tax`
- Output location: `new_paper_data/cognitive_benchmarks/`
- Validation: Compare reported metrics against JSON timestamp `20251222_230849`

.1.4 JSON Log Structure Example

For transparency, here is a simplified example of the JSON structure used for result logging:

```
{
  "scenario": "crypto_tax_compliance",
  "timestamp": "20251222_230849",
  "metrics": {
    "stability_score": 1.0,
    "sec_drift": 0.73,
    "coherence": 1.0,
    "homeostasis_load": 0.32,
    "emotional_calibration": 0.175
  },
  "regulatory_signals": {
    "throttling_activated": true,
    "pathway_conflicts": 2,
    "uncertainty_flags": ["asset_classification", "basis_reconstruction"]
  }
}
```